Bachelor's Thesis

(Course Code: XB_40001)

Psychologists' Views on AI-Supported Mental-Health Tools: An Interview Study

by

Juri Dijkstra

(Student Number: 2653149)

Submitted in partial fulfillment of the requirements
for the degree of
Bachelor of Science
in
Computer Science
at the
Vrije Universiteit Amsterdam

April 27, 2026

Certified by . . . .

Kousar Aslam
First Supervisor

Certified by . . . .

Yulu Wang
Second Reader

Psychologists' Views on AI-Supported Mental-Health Tools: An Interview Study

Juri Dijkstra
Vrije Universiteit Amsterdam
Amsterdam, NL
j.b.dijkstra3@student.vu.nl

ABSTRACT

The growing interest in artificial intelligence as a tool for alleviating pressure on mental health services has prompted questions about how clinicians themselves evaluate these technologies. This study examines how psychologists working in clinical settings reason about the potential role of AI in their work, including which applications they consider appropriate, what concerns they identify, and what conditions they believe must be met before AI can be responsibly integrated into practice. Eleven semi-structured interviews were conducted with psychologists working across clinical and applied settings in five countries.

Three main findings emerged. Participants generally considered AI more appropriate as a support tool than as a replacement for psychologists. Acceptability depended on the type of task, with greater openness toward structured and repeatable functions than toward applications involving direct therapeutic interaction. Several concerns were identified, particularly around privacy, relational limitations, and clinician overreliance. Responsible AI implementation was seen as dependent on human oversight, informed consent, transparent disclosure, careful assessment of task and client suitability, and prior clinical validation.

These findings suggest that AI integration in clinical mental health practice should be approached in a task-specific and carefully bounded way rather than through broad or undifferentiated adoption. They contribute to the emerging evidence base on clinician perspectives by showing that acceptability is not uniform, but depends on the type of clinical task and the risks professionals associate with it. The findings also point to the importance of clear professional guidance on issues such as consent, validation, and oversight as conditions for responsible AI use in psychological care.

Keywords: artificial intelligence, mental health care, psychologists, clinician perspectives, qualitative research, therapeutic relationship, responsible AI integration

1 INTRODUCTION

The Netherlands' mental health services have seen both increased demand and limited capacity in recent years. The Trimbos report "Ggz uit de knel" showed that many people still remain on waiting lists in the Dutch mental healthcare system (GGZ), while workloads for mental health care providers are high and staff shortages continue to grow [3]. These pressures may limit timely access to care for clients and place additional strain on mental health professionals. In response to these perpetual issues, there has been increased interest in potential innovative solutions.

As public attention to Artificial Intelligence (AI) has grown, particularly following the release of ChatGPT and other large language models (LLMs), these technologies are increasingly being discussed as one of the potential tools to alleviate some of the abovementioned pressures on mental health care. Artificial intelligence can be defined as technology designed to perform activities that normally require human intelligence [19]. In mental health care specifically, this includes systems that support clinical reasoning, administrative tasks, and therapeutic processes. In order to better understand where AI might be helpful in alleviating the pressure on the GGZ, it is useful to consider the range of applications currently being explored. Proposed applications include diagnostic support, in which machine-learning models may help clinicians detect, classify, or predict mental health conditions using existing client information, as well as monitoring purposes, for example by processing repeated client data to detect changes over time that may indicate improvement or deterioration [6]. Additionally, AI-supported digital tools have been explored as forms of therapeutic support, including conversational guidance and structured mental health interventions used alongside regular treatment [18]. Each of these applications has the potential to support aspects of clinical work and, in some cases, reduce parts of the clinician's workload.

Despite the potentially beneficial applications, the use of AI in mental health care also raises significant concerns that may complicate its adoption in clinical settings. Multiple reviews have noted that many AI systems rely on data of varying quality [1, 15]. This may contribute to inaccurate or unreliable outputs in clinical contexts, which could in turn have negative consequences for clients. A related concern is the lack of transparency in many AI systems [16]. This makes it difficult for clinicians to understand and critically evaluate the conclusions generated by AI, undermining trust in clinical decision-making. AI systems also create potential privacy risks because they require access to private client data [13, 29]. This raises problems about confidentiality and data security in clinical environments. Finally, biased training data may result in systematically unfair recommendations for specific patient groups [12]. In practice, this could contribute to unequal or discriminatory care for already vulnerable populations. Together, these challenges require not only better technology but also careful professional judgment about whether the potential benefits of a given AI tool outweigh its risks in practice.

Mental health professionals are not passive recipients of AI tools. They are key actors in deciding whether and how these technologies are used in clinical practice and in judging whether a given tool is appropriate for a given client [23, 30]. In psychological care, this professional role includes assessing clients, maintaining the therapeutic alliance, and making final clinical judgments. At the same time, current AI systems still lack capacity in some of the most demanding aspects of mental health work, such as handling complex risk situations, providing safe guidance, and supporting the relational side of treatment [21]. As a result, decisions about whether and how to use AI are not straightforward, and psychologists may need to weigh the potential benefits of a tool against its limitations within their specific clinical context. Understanding how psychologists reason about these tools, including which tasks they consider appropriate for AI support, which they do not, and on what professional grounds, is therefore necessary to identify where AI may responsibly improve practice and where its use becomes problematic.

Much of the existing research on AI in mental health care relies on cross-sectional surveys measuring general perceptions and preparedness [24], thereby providing a broad picture of practitioners' attitudes but limited insight into the reasoning behind them. For example, Cecil et al. [4] found that practitioners' intentions to learn about and use AI varied across application areas. However, even studies that include qualitative components tend to focus on adoption factors such as familiarity, ethical knowledge, and affinity for technology, rather than the professional reasoning that underlies these differences. Such designs can document that variation exists, but they offer limited insight into the professional reasoning behind it. This creates a need for qualitative research on how psychologists judge the potential role of AI in clinical practice.

This study aims to understand how psychologists working in clinical settings evaluate the potential role of AI in their work, including which applications they consider beneficial, which concerns they identify, and under what conditions they believe AI tools could be responsibly integrated into practice.

The study is guided by the following research questions:

- RQ1. What potential roles do psychologists think AI could play in mental health care?
- RQ2. What concerns or risks do psychologists identify when considering the use of AI in clinical practice?
- RQ3. What conditions or safeguards do psychologists identify as necessary for AI tools to be responsibly integrated into clinical practice?

The remainder of this thesis is structured as follows. The background section reviews the existing literature on AI in mental health care, discussing proposed applications and its key challenges. The methods section describes the study design, participant recruitment, and data analysis. The results section presents the main themes that came from the interviews. The discussion interprets these findings in relation to the literature discussed in the background section and reflects on their implications. Finally, the conclusion section summarizes the main conclusions and outlines key considerations for the future development and integration of AI tools in psychological care.

2 BACKGROUND

2.1 Approaches and applications of AI in mental health

AI in mental health care is better understood as a set of different approaches with distinct methods and clinical uses. Recent research has shown that several types of approaches are being explored, such as machine learning, natural language processing, and deep learning [17, 32]. The distinction between these approaches helps to clarify the types of applications and the concerns associated with each. Therefore, it is important not only to make a general description of "AI in mental health", but to go into the different types of systems involved and what they are specifically designed to do.

One important category of AI systems is *supervised machine learning*, which is commonly used for classification and prediction tasks. In supervised learning, algorithms are trained on labeled datasets to learn the relationships between input and output. In mental health settings, these labels may include diagnostic categories, symptom severity, and other clinically defined outcomes. These models may support screening, prediction, and clinical decision support. Reviews in the field show that machine learning and deep learning methods have already been applied to diagnosis-related tasks across several disorders, including depression, anxiety, schizophrenia, bipolar, and post-traumatic stress disorders [14]. These predictive models depend on the quality and validity of the labels on which they are trained, which in turn might impact their accuracy. This may be problematic in mental health settings, where diagnostic categories are contested and may not fully capture clinical heterogeneity. van Oosterzee [25] argues that models trained on such classification systems may reproduce the structure of existing diagnostic systems rather than overcoming their limitations. Their clinical value lies more in supporting assessment and clinical decision-making than in replacing professional judgment.

A second category consists of NLP methods, often including deep learning models, used to analyze unstructured or semi-structured textual data, as opposed to structured inputs. These methods are especially relevant because much clinically meaningful information is expressed through language rather than through neatly structured data. Research in this field has used written sources such as clinical notes, interviews, social media posts, and other textual material to identify patterns associated with mental health conditions [31]. More broadly, systematic reviews show that machine learning and NLP have been applied in mental health research and have potential clinical relevance [17]. Deep learning may be useful for modeling complex relationships in large datasets, but these methods are often less transparent and more dependent on contextually appropriate data.

A third category that has become increasingly well known is *generative AI*, particularly large language models (LLMs). Unlike predictive systems, which mainly try to classify or estimate predefined outcomes, generative AI models produce new text in response to text prompts. This makes them particularly well suited for interactive and language-based functions. Some examples within mental health include conversational agents, self-help or digital intervention tools, and exploratory clinical support applications. A recent systematic review found that LLMs are already being explored for screening, as conversational agents, and a wider set of clinical applications, while also highlighting concerns about accuracy and reliability, interpretability, and overreliance [11]. This represents a technically distinct class of systems because its primary function is not to infer a label from existing data, but instead to generate plausible language dynamically. As a result, generative AI enables an increased level of interaction compared to earlier predictive approaches and may need to be evaluated by partly different criteria.

Distinguishing between these forms of AI is important because they each offer different possibilities and pose different constraints when applied in mental health care. Supervised machine learning is mainly used for prediction and classification, NLP and deep learning may be used for less structured forms of data such as language, and generative AI may offer more interactive applications. These approaches differ in their interpretability, data requirements, reliability, and outputs [11, 17, 25, 32]. Clarifying these categories provides a useful foundation for the remainder of the background. This foundation matters because the discussion that follows does not concern a single uniform technology, but multiple distinct approaches whose implications vary according to their design, application, and intended use.

2.2 Limitations, risks, and human implications

Although AI systems may help with assessment, monitoring, and intervention in mental health care, their use brings limitations that are not only technical but also clinical. Predictive models appear accurate, while also being difficult to interpret, which is an important concern in mental health contexts where understanding model reasoning matters for trust and use in practice [16]. Data-driven systems may also reproduce biases already existing in their datasets [12]. Generative systems may also produce coherent but misleading or incorrect outputs [27]. In mental health settings, these limitations are especially problematic because decisions are often sensitive, context dependent, and closely connected with trust and clinical judgment [16, 27].

One major concern is transparency and interpretability. Many AI systems used in mental health, such as machine learning and deep learning models, function in ways that are not easily understandable to clinicians or patients. As Joyce et al. [16] argues, explainability in mental health cannot be seen as a luxury, since the use of non-transparent systems may lower understandability and make their outputs difficult to justify in practice. This matters for accountability, but might also affect clinical trust. When a model produces a prediction or recommendation without a clear explanation, its practical usefulness may remain limited even if its performance appears good.

A second concern is bias and inequity across populations. AI systems are trained on historical data, and these data may already have preexisting social, diagnostic, or institutional biases. As a consequence, models may reproduce or amplify existing inequities. As Hasanzadeh et al. [12] shows, bias in AI in healthcare can emerge at different stages, including data collection, model development, and implementation, which means that fairness cannot always be assumed, just because a system is data-driven. In mental health care, the expression of symptoms can vary between individuals and groups, while training data cannot fully capture the social and clinical context in which distress occurs.

A third concern focuses on generative AI. Instead of predicting a predefined category or estimating an outcome, these systems generate new text in response to prompts. This may open up possibilities for conversational support, psychoeducation, and assistance for clinicians, but may also introduce several risks. A systematic review by Wang et al. [27] shows that generative AI in mental health is being explored for diagnosis, therapeutic tools, and clinical support, while concerns about hallucinations, reliability, safety, privacy, and ethical oversight still exist. Hipgrave et al. [13] also found that clinicians see potential benefits in generative AI, but are also concerned about confidentiality, inappropriate advice, and the risk of overreliance. Together, these concerns suggest that such systems can seem more reliable or capable than they actually are.

These technical limitations also have broader human implications. Mental healthcare is centered on trust, empathy, shared understanding, and the therapeutic alliance, not only the processing of information. Sauerbrei et al. [22] argues that AI may affect care by changing the experience of empathy, shared decision-making, and trust in the care relationship. Similarly, Weir et al. [28] show that when AI enters mental health care, physicians predict potential benefits and disruptions, especially when tools begin to mediate communication or change care expectations. Additionally, Varghese et al. [26] found that perceptions of AI-driven interventions are shaped by trust, acceptance, and concerns about privacy and appropriateness of AI in mental health care. Together, these findings suggest that the risks of AI in mental health cannot be reduced to technical limitations alone. They also include how systems are understood, whether they are trusted, and how they may change relationships between clients, clinicians, and technologies.

2.3 Practitioner perspectives and research gap

With AI tools becoming more prominent in mental health care, research has shown increasing interest in how professionals understand the role these systems might play in clinical practice. Existing studies suggest that practitioner perspectives are often approached through questions of use, implementation, and perceived benefits and risks. Survey research indicates that mental health professionals recognize potential advantages such as workload support and improved access for clients [5]. Research on AI in mental health services has also emphasized its potential to support clinical decision-making and the importance of implementation conditions for sustained use in practice [2]. Related qualitative work on conversational agents shows that professionals see possible value in these tools, while also identifying important drawbacks, limited capability, and difficulties of integration into care [20]. Taken together, this literature provides valuable insight into the practical and organizational issues surrounding AI adoption. At the same time, it tends to emphasize whether AI can be used more than how professionals reason about where it should or should not be used. This matters because AI systems are not automatically incorporated into practice. Their adoption depends not only on technical performance, but also on whether professionals view them as appropriate for clinical work and aligned with care goals [2].

Therefore, an important gap remains in understanding how psychologists reason about the boundaries and appropriateness of AI use in practice. Existing qualitative research has shown how mental health professionals describe the perceived advantages, disadvantages, and potential future role of AI in care [10]. Other work has focused on how professionals understand the adoption of AI in mental health services and what conditions shape its use [30]. Related research in clinical psychology has focused on barriers and enablers for the implementation of generative AI in practice [7]. Demir et al. [8] similarly explored how mental health professionals reflect on the broader promise and ethical risks of AI in their field. However, what remains less understood is the practical logic through which psychologists evaluate AI differently across administrative, diagnostic, monitoring, and client-facing tasks. This distinction matters because judgments about AI in mental healthcare are likely to involve more than general attitudes toward technology, encompassing task-specific, client-specific, and professional considerations. Whether a given AI application is acceptable in practice cannot be separated from considerations of professional responsibility, therapeutic function, and client vulnerability, all of which shape how psychologists decide what is compatible with good care and what is not. What therefore remains underexplored is how psychologists decide which uses are acceptable and which they regard as incompatible with good care, and which safeguards they consider necessary for responsible integration. This is the gap that the present study addresses.

3 METHODOLOGY

3.1 Research Design

This study employed a qualitative research design using semi-structured interviews. The design was intended to address RQ1 on the potential roles psychologists see for AI, RQ2 on the concerns and risks they identify, and RQ3 on the conditions they consider necessary for responsible integration. A qualitative approach was chosen because the purpose of the study was to examine the reasoning, concerns and contextual judgments underlying the views of the participants. The study examined how psychologists establish professional boundaries about AI use, necessitating participants to articulate their thinking in their own terms, while also allowing the researcher to explore certain distinctions and conditions underpinning these views. Semi-structured interviews are appropriate for this setup, combining a common set of topics with the flexibility to explore examples, clarifications, and new topics raised by participants, consistent with the exploratory aim of the study.

3.2 Participants

The participants were psychologists working in a range of clinical and applied settings, including general mental health care, school psychology, sports psychology, and cognitive rehabilitation. They were based in the Netherlands, Germany, the United States, Spain, and Portugal. The sample was intended to capture the variation in the professional setting, the country, and the level of experience, while maintaining a shared background in psychology. The characteristics of the participants are summarized below, in Table 1.

Table 1: Participant Characteristics (n = 11)

Characteristic	Category	n
Gender	Female	10
	Male	1
Country	The Netherlands	6
	Portugal	2
	Germany	1
	United States	1
	Spain	1
Years of experience	1–3	4
	3–6	5
	6+	2
Organization type	Independent / Personal	3
	Private Organization	4
	Public / Non-Profit Organization	4
Client group	General population	6
	Children / adolescents	3
	Specialized support needs	2

3.2.1 *Selection Criteria.*

Participants were selected to ensure they could provide relevant reflections on the use of AI in psychological practice. Eligible participants had formal training in psychology, were currently working in a clinical or applied setting involving direct contact with clients or patients, had at least one year of professional experience, and were able and willing to participate in an interview in English. Individuals were not included if they did not meet these conditions.

3.2.2 *Sample Size.*

The initial aim of this study was to recruit approximately 15 participants. This target was intended to capture a range of perspectives while keeping transcription and analysis manageable within the scope of a bachelor's thesis. Given the qualitative exploratory design, the goal was not statistical representativeness, but depth and variation in the perspectives of the participants. The final sample consisted of 11 participants, which was considered sufficient to capture diversity across settings, countries, and levels of experience while still allowing for in-depth analysis of individual accounts.

3.2.3 *Recruitment.*

Participants were recruited through convenience and snowball sampling. Initial recruitment was carried out through the researcher's personal and professional network, in which psychologists were approached through mutual contacts. The recruitment was then extended by asking participants to recommend colleagues who might also be willing to participate in the study. Before deciding whether to participate, all potential participants were informed that the study focused on psychologists' perspectives on the use of AI in mental health care.

3.3 Data Collection

3.3.1 *Interview Protocol.*

All interviews were conducted online through video calls on Microsoft Teams. Before each interview, participants were read a brief verbal introduction covering the purpose of the study, the voluntary nature of participation, the use of audio recording, data handling and confidentiality procedures, and their right to withdraw participation or request to exclude their data. Then verbal confirmation of consent was obtained and participants could then ask questions about the study if they had any. Following consent, demographic information was collected, including age, gender, country of practice, professional role, education, years of experience, current practice setting, and treatment population. The interview then lasted roughly 45 to 60 minutes, following the interview guide, while allowing the possibility to explore relevant themes as they arose.

The interview guide was developed based on the research questions and the literature in the background section. Questions on possible applications of AI were included to address RQ1, questions on risks and limitations were included to address RQ2, and questions on safeguards, autonomy, and responsible use were included to address RQ3. The interview guide included questions on participants' familiarity with AI, views on possible applications in diagnostics, monitoring, administrative support, and psychoeducation, as well as risks and limitations. Participants were also asked about the perspective of the client, the therapeutic alliance, and professional autonomy. These topics were included to explore how psychologists reasoned about AI's suitability and the conditions they considered necessary for responsible use in practice. The complete interview guide is included in Appendix A (Interview Guide).

At the end of each interview, participants were thanked and informed about the next steps, including transcription, anonymization, and the deletion of audio recordings after processing.

3.4 Data Processing and Analysis

3.4.1 Transcription. All interview recordings were transcribed using an automated transcription pipeline based on WhisperX[1]. The audio files were processed locally using WhisperX with the large-v3 model. Speaker diarization was used to split the audio and assign basic speaker labels (e.g. `[SPEAKER_00]` and `[SPEAKER_01]`). The transcription was done entirely on a local machine. This ensured data privacy and confidentiality in line with ethical research standards. Generative AI was used to help set up the WhisperX pipeline, no audio or text files were shared.

Following transcription, the diarization labels were manually renamed **P** (Psychologist) or **I** (Interviewer). The transcripts were then lightly cleaned to improve readability. This included only minor edits, such as correcting formatting inconsistencies, removing obvious stammering, and deleting non-meaningful word repetitions, without altering the meaning of participants' responses. The transcripts were then anonymized by removing identifiable information, such as names, ages, company names, regions, and cities. The original transcripts and the cleaned, anonymized transcripts were retained and are available in the accompanying GitHub repository.[2]

Once transcription, verification, anonymization, and cleaning had been completed, all audio files associated with the interviews were permanently deleted.

3.4.2 Data Analysis. The interview data were analyzed in Atlas.ti[3] using an inductive coding approach with an initial topic-based organizing framework. After transcription and anonymization, the cleaned transcripts were read closely and summarized to create an overview of topics and participant responses. Based on these summaries, an initial set of codes was developed.

These codes were then organized into broader code groups corresponding to the main topics covered in the interview guide. This structure was used as an initial organizing framework during analysis, while still allowing for new codes to be added where necessary. All transcripts were subsequently coded by the researcher in Atlas.ti. To illustrate the coding process: when a participant described AI generating draft reports, a code capturing AI-assisted documentation was created and assigned to the pre-existing code group for administrative and practical support. When another participant mentioned automated scheduling, a separate code was created and assigned to the same group. The code groups were established in advance based on the topics covered in the interview guide, while the individual codes within them were developed inductively from what participants said. The complete codebook is available in the accompanying GitHub repository.

The coding process remained iterative throughout the analysis. When new relevant ideas appeared in later transcripts, additional codes were created and assigned to the appropriate code groups. Existing codes were also revised where necessary. In some cases, similar codes were merged into a broader common code, while in other cases codes were split into more specific subcodes when they captured multiple distinct ideas or were mentioned frequently enough to justify further differentiation.

After coding, patterns across the interviews were examined both in terms of recurring views and in terms of meaningful differences between the participants. The final themes were developed from the coded material by comparing patterns across code groups, keeping in mind the specific research questions. A thematic approach was chosen because it allows for systematic comparison across participants on specific topics, and supports conclusions grounded in patterns of convergence and divergence across the sample. A case-by-case approach was considered less suitable, as it would risk interpretive bias where participants did not explicitly explain their reasoning. The analysis focused on how participants reasoned about potential applications of AI, its perceived limitations, and the conditions they considered necessary for responsible AI use. During this process, the researcher also kept informal notes on possible emerging themes and relationships between codes.

3.4.3 Researcher Reflexivity. The researcher's background in psychology and computer science may have influenced the design and interpretation of this study. Prior exposure to psychology shaped an awareness of professional concerns surrounding the use of AI in care. A background in computer science, including machine learning, provided a broader technical understanding of AI systems. This interdisciplinary background was one of the reasons for selecting the topic, as it allowed the researcher to approach the topic from both a psychological and technical perspective.

Prior to conducting the interviews, the researcher anticipated participants would generally be skeptical about the use of AI in psychology, due to having had prior discussions with peers in which AI is often framed as a possible threat to professional roles. Additionally, the researcher anticipated that many participants would primarily associate AI with widely known tools such as ChatGPT, rather than the wider range of applications discussed in the technical literature.

To reduce the influence of these expectations, the interview guide was designed to include questions about both potential benefits and potential concerns, and participants were encouraged to explain their views in their own terms. The researcher was also careful not to steer responses when participants, in response to certain

[1] https://github.com/m-bain/whisperX
[2] https://github.com/juri-bd/thesis-data.git
[3] https://atlasti.com

questions, asked for examples or additional information, and instead emphasized that it was acceptable not to have an answer. During analysis, attention was paid to a range of responses, including positive, negative, and mixed evaluations of AI, as well as to the reasoning participants used to justify them. This reflexive stance was intended to support a more balanced interpretation of the data.

3.5 Ethical Considerations

Ethical approval requirements were assessed using the VU ethical self-check. Based on the outcome of this self-check, no further formal ethical review was required. Participants were informed about the purpose of the study, the voluntary nature of participation, and the handling of their data prior to the interview. Verbal informed consent was obtained before each interview.

3.5.1 Data Storage and Security. Interview recordings and transcripts were stored securely on a local device accessible only to the researcher. Audio files were transcribed locally using the WhisperX pipeline and were not uploaded to external servers. After transcription, manual verification, and anonymization were completed, the audio recordings were permanently deleted. Only the anonymized transcript data were retained for analysis.

4 RESULTS

The results are presented in relation to the three research questions: perceived roles of AI in mental health care (RQ1), concerns and risks associated with AI use (RQ2), and conditions considered necessary for responsible integration into practice (RQ3). Participants differed in their self-rated familiarity with AI, as shown in Table 2. For consistency, the familiarity scores reported here reflect participants' initial self-assessments prior to the interview. In two cases, participants later indicated that they would have rated themselves lower after discussing AI in more detail. More generally, 6 of 11 interviews explicitly reflected uncertainty or limited knowledge about what AI can and cannot do in practice. As one participant noted, "*I know broadly what it does and how something like ChatGPT works, but the depths of it I don't really know*" (Interview 4).

4.1 Perceived roles of AI in mental health care

Addressing RQ1, 7 of 11 interviews described AI as a support tool rather than as a substitute for psychologists. As one participant put it, "*AI for me is like a tool. Not something that you can let manage things.*" (Interview 7). This framework was present in 7 of 11 interviews; participants described AI as something that could assist around care rather than take over clinical responsibility. In practice, 6 of 11 interviews explicitly suggested that the roles considered most appropriate were structured, repetitive and lower-risk, while uses that came closer to replacing clinical judgment or therapeutic interaction were viewed with more caution.

One of the clearest perceived roles was that of administrative and repetitive tasks such as writing reports, summarizing sessions, drafting text, or organizing information. These were seen as particularly suitable for AI because they are relatively standardizable and do not require the kind of contextual judgment that participants reserved for themselves. As one participant explained, "*But the AI makes a really good summary of their sessions. I would be all for it.*" (Interview 1). Administrative support was explicitly mentioned in 6 of 11 interviews, making it one of the more commonly discussed practical use cases.

A second role was found in supporting clinical thinking, for example, brainstorming, generating alternative perspectives, or structuring ideas during assessment or treatment planning. This role involves AI directly interacting with the clinician during ideation. In 5 of 11 interviews, this type of use was framed as exploratory input rather than authoritative guidance. As one participant noted, "*To make an actual diagnosis I wouldn't trust AI, but to brainstorm and get some ideas, I might use it.*" (Interview 2). This role was discussed in 5 of 11 interviews, and in all 5 participants stressed that final interpretation remained the psychologist's responsibility.

Participants also saw potential in AI for monitoring and progress tracking, where AI could help organize repeated observations, summarize changes over time, or present treatment progress more clearly. Four of 11 interviews clearly supported AI for this role, while a further 3 reflected some potential but questioned how much it would add beyond existing practices. In these interviews, participants described monitoring mainly in relation to already structured information, such as repeated observations, questionnaire data, or summarized changes over time. As one participant explained, "*You can put in your prompt what you want documented or monitored and have it produce a graph or summary.*" (Interview 5). Acceptability here was usually tied to the degree to which the information involved was already structured rather than requiring interpretive judgment.

In addition, 5 of 11 interviews described a limited role for AI in psychoeducation or structured client support. The appeal was that AI could add an interactive layer to educational material, complementing the substantial amount of psychoeducational content already available online. At the same time, this role was discussed cautiously, especially because it involves direct client interaction. In these interviews, participants also raised concerns about overreliance and about clients using AI in ways that went beyond the intended scope. These concerns are discussed in more detail in Section 4.2. As one participant explained, "*It could help, as long as it doesn't contradict what I'm saying, because then you have a problem. You have to know what the AI is going to say.*" (Interview 1). This role was discussed more cautiously when it involved direct interaction with clients.

A related but distinct observation concerned the potential appeal of AI for specific clients who find face-to-face interaction difficult. Two participants noted that certain clients, particularly those who are anxious or worried about being judged, might find AI-mediated interaction less socially exposing than direct contact with a therapist. As one participant noted, "*Some clients might prefer AI because you don't see facial expressions*" (Interview 5). Another suggested that for highly anxious clients, "*it can help to start with AI*" as an initial low-pressure point of contact (Interview 10). This was not framed as a general claim that AI makes mental health care more accessible overall. One participant pointed out that clients on long waiting lists may already be turning to AI simply because it is "*cheap and fast*", framing this not as a benefit but as a potential risk of unguided use (Interview 4). Participants who raised this did not describe it as a general claim about accessibility, but as depending on the individual client and situation.

Table 2: Variation in self-rated familiarity with AI

Measure	Range	Values	Note
Self-rated familiarity with AI	2–8	2, 2, 3–4, 4, 5, 6, 6–7, 7, 7, 7, 8	Original familiarity scores were retained for consistency across participants. In two cases, participants later indicated that they would have rated their familiarity lower in retrospect. Some ratings were given as ranges rather than single values.

4.2 Concerns and risks regarding AI use in clinical practice

In relation to RQ2, one of the most prominent concerns that emerged during the interviews was related to privacy and data security. Nine of 11 interviews raised caution about whether client information could be stored, processed, or protected safely when AI was used. They found this problem especially alarming due to the sensitivity of the data being shared in clinical practice, such as sexuality, medical information or other personal information. If such information gets exposed, it could have very negative effects on their personal and professional lives. As one participant explained, "*I have so many concerns... mostly because the information is so sensitive and we're already doing so much to make sure that people don't even have access, let alone that it gets on a data server.*" (Interview 8). Privacy and data protection concerns were raised in 9 of 11 interviews.

Eight of 11 interviews explicitly raised the concern that the relational and nonverbal aspects of care cannot be replicated by AI. Some of these include body language, empathy, and the flexibility of the therapist to respond to what happens in the room. All of these were described as being essential to providing good care for clients. A specific example was the concern about losing access to transference and countertransference, the emotional responses that arise between therapist and client, which several participants described as clinically meaningful information in itself. As one participant stated, "*Limitations like transference or countertransference in therapy. When a patient makes you feel a certain way, or reacts in a specific way, and feelings arise. ChatGPT or Gemini cannot feel that, and I think that's very valuable in therapy, especially in psychotherapy.*" (Interview 6). This concern was raised in 8 of 11 interviews.

Another concern was that AI could produce misinformation, hallucinations, or unsafe outputs. This concern was raised explicitly in 6 of 11 interviews, particularly in relation to client-facing use. Participants noted that AI might give information that is incorrect or inappropriate for a specific client's situation. Using AI with vulnerable clients was seen as especially risky. In general misinformation is worrying, but it becomes especially dangerous with clients who are in crisis, experiencing psychosis, or prone to suicidality, since AI's responses may cause direct harm to the client. As one participant explained, "*So I would be careful with AI because you don't want clients to get wrong information. Even if it might not be wrong, it's not checked. You don't really see what they are getting.*" (Interview 4). A related concern, raised in 4 of 11 interviews, was that AI may validate rather than challenge clients, whereas therapy often requires a more reflective and sometimes confronting role. As one participant explained, "*these beliefs that are seen on an AI platform or in digital sources, they are validating primary beliefs... So it's this tendency to validate beliefs that may not be the best for the person.*" (Interview 11).

A related concern was that AI may rely too heavily on general patterns while not taking into account the individual client and their context. This concern was raised explicitly in 5 of 11 interviews. Participants pointed out that similar behavior can mean very different things depending on a person's background, history, and situation, and that AI systems won't be able to capture these kinds of nuances reliably. This problem was most clearly expressed in a specific example given by a participant: "*For example, if I work with an athlete who is really strict about sleeping times and eating habits and is very rigorous about that, there might be a chance AI says this is autism, since there are some symptoms in there. But those are also things you sometimes need in elite sports to perform well. In the sports context it's not necessarily a symptom of autism, while in another context it could be.*" (Interview 2). Participants used examples like this to explain why they felt AI could not reliably account for individual context.

Six of 11 interviews also noted that clients are already using AI outside of therapy sessions, with the most concrete concern being inconsistency between AI-generated information and the therapist's treatment approach. If a client receives information from an AI tool that contradicts the treatment plan, this could negatively affect the therapeutic relationship. As one participant put it: "*If clients get different information from AI than from me, that could cause tension.*" (Interview 4). Another described the risk more plainly, suggesting that when AI contradicts the therapist, and clients trust the AI more than the human clinician, this can harm the therapeutic relationship (Interview 10). Related to this was the concern for clients to become overreliant on AI and the effect it may have on emotion regulation and coping. Four of 11 interviews noted that clients might become too comfortable interacting with AI, therefore taking away the opportunity to have real-world help seeking or interpersonal connections. As one participant observed, "*Students are unlearning how to interact with others in person,*" adding that clients could become "*really connected to a bot*" in ways that make human interaction feel harder over time (Interview 5). These concerns were not consistently raised across all interviews, but participants who raised them described unease about clients substituting AI interaction for human connection or therapeutic work.

Four of 11 interviews raised additional concerns about the impact of AI on the therapeutic alliance. This came up less frequently than other concerns, and most participants did not describe AI as taking over the therapeutic role of the psychologist. However, 4 of 11 interviews noted that more limited uses of AI could already affect how clients feel in the therapeutic relationship. This extended beyond client-facing tools to background uses such as session recording or AI-generated summaries. As a consequence, participants suggested that clients may feel less heard or become more reserved, since some clients might not be comfortable speaking when they feel like something else is observing them besides their clinician. As one participant stated, "*The first sessions are very important to establish a therapeutic alliance. So if you introduce a tool in the first session, I think it disrupts that.*" (Interview 7). These concerns were especially prominent for client-facing uses, but extended to any AI use that clients were aware of during sessions.

A further concern, distinct from the expectation of human oversight discussed in Section 4.3, was that clinicians themselves might become overly reliant on AI in ways that reduce their clinical attention. Where Section 4.3 addresses what responsible use requires, this concern addresses what happens when those requirements are not met. Five of 11 interviews raised concerns that clinicians might begin to trust AI output too quickly, check it less carefully, and gradually let it stand in for their own judgment. As one participant noted, "*people could get lazy and rely on it too much without making sure the information or interpretation is correct*" (Interview 5). Another framed this in more directly clinical terms: "*if you trust AI too much, I think you don't see the clients, how he or she really is*" (Interview 9). The concern here was that clinicians might shift attention away from the person in front of them and toward the AI output.

Finally, 6 of 11 interviews pointed to uneven organizational handling of AI use, while 4 of 11 interviews explicitly suggested that disclosure practices were inconsistent across practitioners or settings. Evidence from the interviews suggested that clients do not always know when AI is being used in their treatment. As one participant noted, "*Usually they don't know, and it depends on the professional*" (Interview 6), implying that disclosure was inconsistent across practitioners rather than a settled standard. This unevenness was mirrored at the organizational level. Six interviews described workplaces where AI governance was partial, still developing, or handled differently across settings. One participant reported that their organization had issued only a basic portal message saying staff could use Copilot but should not enter private client information, and said that beyond this the organization was "*not supporting except of the message on the portal for the people who work there*" (Interview 9). Another described their organization as still working out which tools were preferred and why (Interview 4). Across these interviews, participants described AI use as taking place under uneven organizational conditions, with key decisions often left to individual clinicians.

4.3 Conditions for responsible integration of AI into practice

Addressing RQ3, participants described responsible AI integration as dependent on a combination of conditions rather than a single safeguard, including professional oversight, privacy protection, transparency toward clients, and the appropriateness of the task and client group involved.

A condition raised in 9 of 11 interviews was that professional responsibility should always remain with the psychologist. Participants stressed that even when AI could help with some of the applications mentioned in Section 4.1, the final clinical judgment should always be borne by the psychologist and may never be delegated to an AI system. The psychologist should remain responsible for how the AI output was interpreted, for its accuracy, and its application in practice. As one participant explained, "*As long as I can check it, I feel in control and responsible. I don't think it's the AI's responsibility.*" (Interview 3).

Closely connected to this, but distinct from it, was the expectation that AI-generated output should always be actively verified by the psychologist. Where the previous condition concerns who holds final authority, this one concerns the practical act of review: participants did not treat AI output as something that could simply be passed on to clients or incorporated into clinical decisions without being checked first. This expectation was raised in 8 of 11 interviews. As one participant explained, "*I absolutely think it's still your responsibility to check what the AI is doing, whether you stand behind it and whether you think it's accurate.*" (Interview 8). This expectation was closely linked to concerns about accountability and the possibility of inaccurate or misleading outputs discussed in Section 4.2.

Another condition concerned privacy and data security. Participants often described strong safeguards around confidentiality as one of the most important prerequisites for acceptable AI use in practice. As one participant explained, "*What I need is to know that the data I enter is safe, just in case I accidentally enter some personal information.*" (Interview 6). Given the highly sensitive data, as described in Section 4.2, this condition was raised in 9 of 11 interviews and was often described as non-negotiable.

Participants also mentioned that the type of task influences the acceptability of AI use. AI was generally considered more appropriate for lower-risk, structured tasks than those requiring complex interpretation, sensitivity to crisis, or empathy. Responsible integration was therefore often described as restricting AI to tasks where the stakes of error were lower. As one participant explained, "*for sure to try it for a while and to really give specific tasks to it. Like all the administration stuff, all the behind the scenes stuff, I would trust it more to outsource than the responsibility over a living human being.*" (Interview 10). This boundary between support functions and tasks requiring human judgment was raised in 9 of 11 interviews when participants described what responsible AI use would look like in practice.

Clinically acceptable AI use was described in 7 of 11 interviews as depending on the specific client and context. Participants approached this similarly to how they described selecting other clinical tools or methods: suitability was not something determined by a fixed rule but by the individual psychologist's judgment about what a particular client was likely to respond well to. As one participant noted, "*I think it depends on the client, whether that would be to your benefit all the time or not. ... For example, if you have OCD or anxiety that could go the wrong way, or an addiction pattern.*" (Interview 3). Age was mentioned as one relevant factor in 6 of 11 interviews. In these interviews, participants generally expected younger clients to be more open to AI, with one estimating that clients under 25 or 30 would rate around a seven or eight in openness while those over 40 or 50 would be much lower, because older clients were seen as less comfortable with how AI works (Interview 5). That said, participants did not treat this as a universal rule. Some did note that older clients might be more willing to use it when framing was done correctly by the psychologist, and that age mattered less than the actual clinical problem and AI application use. In practice, participants suggested that determining suitability would require looking at factors such as vulnerability, age, the nature of the clinical problem, and the stage of treatment. Once these have been considered, it is then the task for the clinician to discuss the appropriateness of the tool with the client to determine eventual application.

Six of 11 interviews emphasized the importance of consent, disclosure, or client awareness as conditions for acceptable AI use. In one setting, clients were explicitly informed before responding that suggestions were generated by AI under the therapist's supervision, with consent obtained as a procedural step rather than simply as part of an explanation (Interview 7). In other interviews this was described more informally, as a discussion in which the therapist introduced the tool, explained its purpose, and checked whether the client was willing to proceed. As one participant stated, "*So I think if there is a way of using AI during the sessions as a tool, people could accept it because it would already be validated by the therapist. So the therapist would choose this tool to use, as we use Google Teams or other things. But first, I think it has to make sense to the therapist so we could introduce this tool like any other tool.*" (Interview 11). The variation in how open participants believed their clients to be to AI use is shown in Table 3. The wide range of scores (3 to 9) does not reflect uniform client attitudes but differences in client group, type of AI use, and how clearly its use was communicated. Participants described consent and disclosure as necessary conditions, but noted that perceived acceptability also depended on the clinical context.

Four interviews indicated that explainability was a prerequisite for ethically acceptable AI use in practice. Before trusting AI output enough to use it clinically, participants wanted to understand where the system gets its information and how it reaches its conclusions. As one participant asked directly, "*Where does the information come from and how did it reach that conclusion?*" (Interview 4). Another similarly stated that clinicians would need to know "*where it gets its information from*" before relying on its output (Interview 2). Participants who raised this condition described it as necessary before they would feel comfortable relying on AI output clinically.

Finally, two participants explicitly described prior testing and validation as a condition for acceptable use. Before an AI tool could responsibly be introduced into clinical practice, participants wanted some assurance that it had been evaluated and that its outputs could be trusted. As one participant stated, "*I think it has to be pre-approved, with a lot of rigorous testing beforehand*" (Interview 3). The same participant also suggested that ongoing monitoring would be necessary, describing the value of "*a second instance constantly checking the responses, a feedback loop of sorts.*". The two participants who raised this condition described it not only in terms of prior scrutiny, but also as requiring continued monitoring rather than a one-time approval. One participant emphasized formal pre-approval and rigorous testing beforehand (Interview 3), while the other stressed the need to understand how the system works before relying on its output (Interview 11).

5 DISCUSSION

How psychologists think about AI in their work appears to be more nuanced than some of the existing literature suggests. This study aimed to answer this question through eleven semi-structured interviews with psychologists. These interviews focused on the applications they considered appropriate, the concerns they had, and the conditions that need to be met before feeling confident in using AI responsibly in clinical practice. Three main findings emerged. (i) Participants were not outright opposed to AI; they often instead drew distinctions between tasks they considered suitable for AI support and tasks that were not, (ii) Participants were generally less accepting of AI integration the closer it came to direct therapeutic interaction, and (iii) Privacy and the relational dimensions of therapy were among the most prominent concerns, but participants also raised issues such as clinician overreliance and the absence of clear organizational guidance. In their view, responsible integration was not a matter of simply meeting a single condition. It involved a combination of intersecting requirements, including oversight, explainability, consent, task and client suitability, and prior validation, which participants often treated as cumulative rather than interchangeable. The sections that follow interpret these findings in relation to the existing literature and reflect on what they mean for clinical practice and future research.

5.1 AI as a supportive tool: the task-appropriateness gradient

The finding that participants generally framed AI as a supportive rather than substitutive tool aligns with existing literature emphasizing that AI's clinical value lies in assisting professional judgment rather than replacing it. As van Oosterzee [25] argues in the context of supervised machine learning, predictive models are better understood as decision-support tools than as autonomous diagnostic systems. Participants in this study extended that logic across a broader range of applications, applying it not only to diagnostic or predictive tasks but also to administrative work, clinical brainstorming, and monitoring. What the results add beyond this general framing is a more specific pattern: acceptability was not uniform across tasks but followed a gradient. Participants were most accepting of AI for structured, administrative, and lower-risk tasks, relatively open to AI supporting clinical thinking and monitoring, and noticeably more cautious about any AI use involving direct client interaction. This gradient roughly reflects a distinction between tasks where AI processes existing structured information and tasks where it generates or mediates clinical meaning. The former were broadly accepted, the latter were treated with significantly more caution.

This pattern is consistent with findings from Cross et al. [5], who found that mental health professionals recognized workload support as one of the clearest potential benefits of AI, and with Cecil et al. [4], who found that practitioners' intentions to use AI varied across application areas. However, neither of those studies fully explains the underlying logic driving that variation. The findings from this study suggest that differences in acceptability were not explained by familiarity alone, but were also shaped by how closely a task approached the relational and interpretive core of therapeutic work. This is a more specific formulation than is typically made in the literature reviewed here, and it suggests that task type may function not only as a moderating factor in AI adoption, but also as an important professional boundary.

Table 3: Variation in perceived client openness to AI use in treatment

Measure	Range	Values	Note
Perceived client openness to AI use in treatment	3–9	3, 3–4, 3–4, 4, 5–7, 6, 6, 6, 7, 7, 8–9	Ratings varied and were often described as depending on the type of AI use, the client group involved, and the extent to which the use of AI would be explained transparently. Some ratings were given as ranges rather than single values.

The finding also connects to the background distinction between AI approaches. Supervised machine learning and structured monitoring tools sit comfortably within the administrative and data-handling end of the gradient. Generative AI, with its capacity for interactive and language-based functions, sits closer to the client-facing end, which may partly explain why participants were more cautious about psychoeducation and conversational applications than about summarization or progress tracking. Guo et al. [11] noted similar concerns about overreliance and reliability in generative AI applications, and this study suggests that these concerns are not evenly distributed across all use cases, but rather cluster around applications that are closer to therapeutic interaction.

5.2 Privacy, relational care, and the limits of AI in therapy

Privacy was one of the most prominent concerns raised across the interviews. This is consistent with existing literature. Hipgrave et al. [13] found that clinicians identified confidentiality as a central concern when considering generative AI, and Yoo et al. [29], in a study examining clinician perspectives on AI-derived insights from patients' social media activity, similarly raised privacy concerns about AI systems handling sensitive client data. What the present findings add is a more specific account of why privacy concerns may feel especially acute in psychological care. Participants often linked privacy not only to data protection in a general sense, but also to the particularly sensitive nature of therapeutic information itself. Psychological data can reveal highly sensitive information about a person's mental health, identity, and personal struggles. Participants pointed out that this information getting exposed could have serious consequences in the personal and professional lives of clients. The current study therefore suggests that confidentiality is not only important for compliance, but is also closely tied to the trust on which participants understood therapeutic work to depend.

The second major concern, appearing across many interviews, was that participants viewed the relational and nonverbal dimensions of therapy as beyond what AI could adequately replicate. Participants described body language, tone, empathy, timing, and the therapeutic use of transference and countertransference as central to good care and, in their view, not accessible to current AI systems in a clinically meaningful way. This finding connects to Sauerbrei et al. [22], who discuss how AI may affect care by altering the experience of empathy, shared decision-making, and trust in the therapeutic relationship, and to Weir et al. [28], who found that clinicians anticipated disruptions to therapeutic communication when AI began to mediate the care relationship. The present findings are consistent with both. What participants in this study added was a more specific account of what they saw as difficult for AI to capture: not empathy in general, but body language, tone, timing, nonverbal cues, and the interactional dynamics between therapist and client, including transference and countertransference.

Taken together, these two concerns point to a shared logic in how participants assessed AI's limits in clinical work. Privacy concerns reflect the sensitivity of what is shared within the therapeutic relationship; relational concerns reflect participants' view that the process through which that sharing happens is inherently human. Both were often treated not simply as practical preferences, but as important conditions for therapy. This may help explain the pattern described in Section 5.1: the closer a task comes to the relational core of therapy, the more directly it implicates both concerns, and the less acceptable AI involvement becomes.

5.3 Clinician overreliance and organizational gaps as underexplored risks

Two concerns that emerged from the interviews receive relatively little attention in the existing literature reviewed for this study. The first was that some participants worried clinicians themselves might become overly reliant on AI in ways that reduce clinical attention. Participants worried that clinicians might trust AI output too quickly, check it less carefully, and allow it to shape their judgment in ways that draw attention away from the person in front of them and toward the AI output. This concern is distinct from the checking and oversight conditions described in Section 4.3, which address what responsible use requires. Here, the concern is what may happen when those requirements are not met in practice.

In the literature reviewed for this study, overreliance was more often discussed in relation to client-facing applications than in relation to clinicians' own use of AI. Hipgrave et al. [13] identify overreliance as a concern in the context of clients engaging with generative AI tools, and Guo et al. [11] similarly flag it in conversational AI applications directed at users. The present findings suggest that a related risk may also apply on the clinician side. Whether this dimension has received insufficient attention in the broader literature is difficult to assess from the studies reviewed here, but it did not feature prominently in the sources consulted, and several participants raised it in relation to professional practice. This points to clinician-side overreliance as a direction that future research would benefit from examining more directly.

The second concern was the absence of clear organizational guidance on AI use and the inconsistency of disclosure practices across settings. Participants described workplaces with only minimal policies, partial pilots, and unresolved questions about which tools were appropriate and why. Four of the eleven participants noted that clients did not always know when AI was being used in their treatment, with disclosure depending on the individual practitioner rather than on any shared standard. Auf et al. [2] emphasize that implementation conditions matter for sustained and responsible AI use in mental health services, and Zhang et al. [30] similarly identify organizational factors as shaping how professionals adopt AI in practice. Where those accounts address implementation conditions at a framework level, the present findings give a more concrete picture of what the absence of those conditions looks like from the perspective of individual clinicians: practitioners often navigating these decisions with limited institutional support, clear ethical guidance, or shared disclosure norms.

This has practical implications. When organizational guidance is lacking, individual psychologists are left to decide for themselves how AI should be used, while possibly lacking the technical knowledge, time, or institutional support needed to make well-informed decisions about these tools. Participants' accounts suggest that this can contribute to inconsistent practice across settings, raising questions about whether individual professional responsibility alone is a sufficient basis for reliable AI use in clinical contexts. Whether and how institutions are currently addressing this gap is an empirical question this study was not designed to answer, but the findings point to organizational and regulatory action as a necessary complement to individual professional judgment.

5.4 Conditions for responsible integration

The conditions participants described for responsible AI use were not a simple checklist. Professional responsibility remaining with the psychologist, active verification of AI output, privacy safeguards, task and client suitability, consent and disclosure, explainability, and prior validation were all identified as important, and participants generally treated them as cumulative. What this means in practice is that a single condition being met was rarely enough on its own: a tool might handle privacy adequately but still be considered unacceptable if the task was inappropriate, or if outputs could not be traced back to an understandable source. Responsible use, in participants' accounts, depended on several of these conditions being satisfied together.

This finding is broadly consistent with Auf et al. [2], who emphasize that sustained responsible use of AI in mental health services depends on implementation conditions rather than technical performance alone, and with Zhang et al. [30], who identify professional and organizational factors as shaping whether AI adoption is appropriate in clinical contexts. The present findings extend this by specifying what those conditions look like from the perspective of practicing psychologists, grounding them in the concrete reasoning participants used rather than deriving them from implementation frameworks or policy documents.

Two of the conditions identified deserve particular attention. The first is explainability. In 4 of 11 interviews, explainability was treated not merely as a desirable feature, but as a prerequisite: understanding where a system gets its information and how it reaches its conclusions was described as necessary before AI output could be meaningfully checked or responsibly used. This connects directly to Joyce et al. [16], who argue that explainability in mental health AI cannot be treated as a luxury since non-transparent systems undermine the accountability and trust that clinical use requires. Findings from this study support that argument from the practitioner side, suggesting that explainability is not only a technical or ethical requirement, but a practical condition that shapes whether oversight is even possible.

A further condition that emerged was contextual and population fit. Participants suggested that responsible AI use depended on whether a given tool was appropriate for the specific client, setting, and clinical context in which it was used. Some pointed out that psychological assessment is shaped by population norms and cultural context, meaning that outputs cannot simply be assumed to apply across countries or client groups. Others noted that certain presentations, such as complex comorbidity or particular personality structures, could make AI use less appropriate or more difficult to introduce responsibly. Participants treated suitability as something that had to be evaluated in relation to the clinical situation at hand.

Taken together, the conditions participants described suggest that responsible AI integration is not primarily a matter of individual willingness but also of structural readiness. Several of the conditions discussed by participants, particularly organizational disclosure standards and regulatory oversight, are not things individual clinicians can arrange on their own. Some participants also linked responsible use to prior validation. Participants' accounts suggested that where these structural conditions are missing, practitioners are largely left to make do without the institutional backing that responsible use, as they described it, would seem to require.

5.5 Limitations

Several limitations of this study should be acknowledged when interpreting its findings. Participants were recruited through convenience and snowball sampling; this resulted in a non-random sample where psychologists who are more interested or open to AI may be overrepresented. Some participants mentioned that colleagues they approached declined participation because they held negative views about AI use in mental health care. This suggests that more negative views on AI may be underrepresented in the sample. To reduce this risk, efforts were made to recruit as many participants as possible through available networks and follow-up outreach. However, within the time constraints of the thesis and with limited response rates, the final sample remained small and non-random. The findings should therefore be interpreted with caution, and future research would benefit from broader and more systematic recruitment strategies.

The sample was relatively small at eleven participants, with six of the eleven participants based in the Netherlands. Although the sample included some variation in national and professional context, the relatively small sample means that some relevant perspectives, concerns, or conditions may not have been captured. In addition, the concentration of participants in one national setting means that the findings should not be treated as broadly representative of all psychologists working in mental health care. Future research would benefit from larger samples and more balanced recruitment across national and clinical contexts.

For many of the participants, English was not their first language. The decision to conduct interviews in English had two primary reasons. One reason was that the study was conducted as part of an English-language degree program, meaning that the transcripts and analysis ultimately had to be in English. Another reason was to have all transcripts in one shared language, which made comparison across interviews more consistent. Where needed, follow-up questions were used to clarify participants' meaning during the interviews. At the same time, conducting the interviews in English may have limited how precisely some participants could express their views, and some nuance in their answers may therefore have been lost. Future research could address this by conducting interviews in participants' native languages or by comparing findings across different language contexts.

Participants' self-reported familiarity with AI was mixed and should be interpreted with caution. Two participants noted during the interview that they would have rated their familiarity lower at the end of the interview than the score they gave at the start, suggesting that they might have initially overestimated their own familiarity levels. To reduce the impact of this inconsistency, the original ratings were retained for consistency across participants, but they were interpreted cautiously and considered alongside the interview content rather than as precise objective measures of familiarity. Future work would benefit from including psychologists with more direct hands-on experience of AI tools and from using more concrete measures of familiarity or use in addition to self-report.

Interpretive judgment was used in many stages of this study, including the background research, the compilation of the interview questions, the use of follow-up prompts, and the coding of the material. All interviews were conducted, coded, and analyzed by a single researcher. This may have increased consistency in procedure across interviews, but no full independent verification or inter-rater reliability checks were carried out, apart from supervision during the coding of the first interviews. Researcher influence and bias may therefore have affected the design and interpretation of the findings. Future research would benefit from involving additional coders, independent verification, or other forms of analytic triangulation.

5.6 Implications for practice and future research

The findings of this study have several implications that should be considered. For psychologists, the task appropriateness gradient identified in this study offers a practical guideline for determining which types of AI applications may be appropriate to implement. Administrative tasks, brainstorming and more structured monitoring generally seem more accepted, while client-facing applications and those closer to direct therapeutic interaction require more caution. Before implementing applications, psychologists may need to consider whether a tool is appropriate for the task, the client population, and the clinical context involved. In practice, AI seems easier to justify in supportive and lower-stakes functions, and more difficult to justify in applications that come closer to the relational core of therapy.

For organizations and institutions, the findings suggest that individual professional judgment is not enough to be able to make a consistent use of AI in clinical settings. The inconsistency in disclosure practices and the minimal organizational guidance described by participants point to a need for clearer institutional policies. These should cover the types of tools approved for use and under what conditions they may be used and how to use them. Without such guidance, the burden falls on individual clinicians who may already be overworked and may lack the technical expertise or institutional support to navigate these decisions reliably.

For policymakers and regulators, the findings suggest a need for clearer legal and professional rules on how AI may be used in clinical care. Participants' accounts indicate that not all forms of AI use should be treated as equally acceptable: some applications may need to be restricted, some may only be appropriate in limited supportive roles, and uses that come closer to direct therapeutic interaction may require especially strong safeguards. The findings also suggest the need for clearer guidance that states that the responsibility for clinical decisions remains with the psychologist, even when AI is used as a support tool.

For future research, several directions appear worthwhile. One is further investigation of clinician-side overreliance. Although this was not among the most dominant themes in the interviews, some participants did raise concerns that clinicians might trust AI output too quickly or check it less carefully, suggesting a need for more direct empirical work on how AI use may affect clinical attention and judgment over time. Another is research on institutional governance. Participants' accounts pointed to uneven organizational guidance and inconsistent disclosure practices, indicating a need for studies examining how mental health institutions currently regulate AI use and what kinds of support or policy structures may help promote more consistent and ethical practice. Future qualitative work could also benefit from broader sampling across professional orientations, national contexts, and levels of AI experience, since the present sample was limited in size and the reasoning of practitioners with more direct experience of clinical AI tools may differ from those with little or no hands-on use.

6 CONCLUSION

This study examined how psychologists reason about AI in clinical practice, drawing on semi-structured interviews with eleven participants across five countries. It explored what roles they considered AI suited for, what concerns they associated with its use, and what conditions they believed would need to be met before responsible integration was possible.

Participants saw AI primarily as supportive for specific tasks, and acceptability followed a consistent gradient. Tasks that were repetitive and more standardized, such as administrative work, clinical brainstorming, and structured monitoring, were generally considered more appropriate. Client-facing applications and uses that came closer to clinical judgment or the relational dynamics between therapist and client were viewed with considerably more caution and were largely rejected when they approached the interpretive or relational core of therapeutic work.

The most widely shared concerns were privacy and the relational foundations of therapy. Psychological data was described as uniquely sensitive, and the therapeutic relationship was seen as dependent on forms of human presence, attunement, and emotional responsiveness that participants viewed as inaccessible to current AI systems. Participants also identified risks that receive less attention in the existing literature, including clinician overreliance and the absence of clear organizational guidance, with disclosure practices varying considerably across settings.

On the conditions for responsible integration, participants described a combination of overlapping requirements including professional oversight, active verification of AI output, privacy safeguards, task and client suitability, consent and disclosure, explainability, and prior validation. These were generally treated as cumulative. Several cannot be met by individual clinicians alone, pointing to the need for coordinated action at the organizational and regulatory level.

The primary contribution of this study is an empirically grounded account of how practicing psychologists reason about AI, organized around a task-appropriateness gradient and a set of conditions that go beyond what existing implementation frameworks typically specify. These findings suggest that psychologists hold nuanced, context-dependent views rather than fixed positions for or against AI. Tools developed without engaging with this reasoning may struggle to gain clinical acceptance regardless of technical performance.

Several directions for future research follow from these findings. Clinician-side overreliance received little attention in the existing literature, but was raised consistently enough to warrant a dedicated empirical investigation. Studies examining how institutions currently govern AI use in mental health settings would help clarify the organizational gaps described by the participants. Future qualitative work would also benefit from broader sampling across national contexts and levels of AI experience, since the reasoning of practitioners with direct hands-on experience may differ meaningfully from those with little or no such experience.

Responsible integration will require AI applications that offer clear clinical or practical value, alongside stronger organizational guidance, clearer disclosure standards, and a continued commitment to keeping the core of therapy under the clinician's control.

7 ACKNOWLEDGMENTS

The author sincerely thanks Kousar Aslam for her valuable guidance, thoughtful feedback, and continued support throughout this thesis. The author also warmly congratulates her on the upcoming arrival of her baby and wishes her all the best for the coming period and her maternity leave. The author also sincerely thanks Yulu Wang for serving as second reader and for her time, feedback, and involvement in the evaluation of this thesis.

The author acknowledges the use of generative AI tools during the preparation of this thesis. In line with the European Commission's guidelines on responsible use of generative AI in research [9], these tools were used only in a supportive capacity, primarily to assist with brainstorming and to test whether the logic and phrasing of sections were clear. They did not make decisions about the inclusion of sources and the AI-generated text was not copied directly into the thesis. Responsibility for the final wording, the selection and verification of sources, the interpretation of the literature, and all academic judgments rests entirely with the author.

REFERENCES

[1] Maisam Ali, Shahid Ali, Qaiser Abbas, Zeeshan Abbas, and Seung Won Lee. 2025. Artificial intelligence for mental health: A narrative review of applications, challenges, and future directions in digital health. *Digital Health* 11 (2025), 20552076251395548. https://doi.org/10.1177/20552076251395548

[2] Hassan Auf, Petra Svedberg, Jens Nygren, Monika Nair, and Lina E. Lundgren. 2025. The Use of Artificial Intelligence in Mental Health Services to Support Decision-Making: Scoping Review. *Journal of Medical Internet Research* 27 (2025), e63548. https://doi.org/10.2196/63548

[3] Jenny Boumans, Hans Kroon, and Bert van der Hoek. 2023. *Ggz uit de knel: Verkenning van mogelijkheden en valkuilen van het actuele hervormingsnarratief.* Technical Report. Trimbos-instituut, Utrecht, The Netherlands. https://www.trimbos.nl/wp-content/uploads/2025/10/AF2059-Ggz-uit-de-knel.pdf

[4] Julia Cecil, Anne-Kathrin Kleine, Eva Lermer, and Susanne Gaube. 2025. Mental health practitioners' perceptions and adoption intentions of AI-enabled technologies: an international mixed-methods study. *BMC Health Services Research* 25, 1 (2025), 556. https://doi.org/10.1186/s12913-025-12715-8

[5] Sabrina Cross, Ian Bell, Jennifer Nicholas, Louise Valentine, Sarah Mangelsdorf, Stuart Baker, Nick Titov, and Mario Alvarez-Jimenez. 2024. Use of AI in Mental Health Care: Community and Mental Health Professionals Survey. *JMIR Mental Health* 11 (2024), e60589. https://doi.org/10.2196/60589

[6] Pablo Cruz-Gonzalez, Aaron Wan-Jia He, Elly PoPo Lam, Ingrid Man Ching Ng, Mandy Wingman Li, Rangchun Hou, Jackie Ngai-Man Chan, Yuvraj Sahni, Nestor Vinas Guasch, Tiev Miller, Benson Wui-Man Lau, and Dalinda Isabel Sánchez Vidaña. 2025. Artificial intelligence in mental health care: a systematic review of diagnosis, monitoring, and intervention applications. *Psychological Medicine* 55 (2025), e18. https://doi.org/10.1017/S0033291724003295

[7] Daniel de la Fuente Tambo, Sara Iglesias Moreno, and Manuel Armayones Ruiz. 2025. Barriers and enablers for generative artificial intelligence in clinical psychology: a qualitative study based on the COM-B and theoretical domains framework (TDF) models. *BMC Psychology* 13, 1 (2025), 1181. https://doi.org/10.1186/s40359-025-03500-7

[8] Mehmet Demir, Fuat Aydoğdu, Mine Begümhan Alabay, Hatice Kübra Yaşar, and Aras Bozkurt. 2026. More than an algorithm: mental health professionals confront the promise and ethical perils of artificial intelligence. *Universal Access in the Information Society* 25, 1 (2026), 24. https://doi.org/10.1007/s10209-025-01297-1

[9] European Commission. 2025. *Living guidelines on the responsible use of generative AI in research.* ERA Forum Stakeholders' document. European Commission, Directorate-General for Research and Innovation, Brussels, Belgium. https://research-and-innovation.ec.europa.eu/document/download/2b6cf7e5-36ac-41cb-aab5-0d32050143dc_en?filename=ec_rtd_ai-guidelines.pdf Second Version, April 2025.

[10] Mücahit Gültekin and Meryem Şahin. 2024. The use of artificial intelligence in mental health services in Turkey: What do mental health professionals think? *Cyberpsychology: Journal of Psychosocial Research on Cyberspace* 18, 1 (2024), Article 6. https://doi.org/10.5817/CP2024-1-6

[11] Zhijun Guo, Alvina Lai, Johan H. Thygesen, Joseph Farrington, Thomas Keen, and Kezhi Li. 2024. Large Language Models for Mental Health Applications: Systematic Review. *JMIR Mental Health* 11, 1 (2024), e57400. https://doi.org/10.2196/57400

[12] Fereshteh Hasanzadeh, Colin B. Josephson, Gabriella Waters, Demilade Adedinsewo, Zahra Azizi, and James A. White. 2025. Bias recognition and mitigation strategies in artificial intelligence healthcare applications. *npj Digital Medicine* 8 (2025), 154. https://doi.org/10.1038/s41746-025-01503-7

[13] Lyndsey Hipgrave, Jessie Goldie, Simon Dennis, and Amanda Coleman. 2025. Balancing risks and benefits: clinicians' perspectives on the use of generative AI chatbots in mental healthcare. *Frontiers in Digital Health* 7 (2025), 1606291. https://doi.org/10.3389/fdgth.2025.1606291

[14] Ngumimi Karen Iyortsuun, Soo-Hyung Kim, Min Jhon, Hyung-Jeong Yang, and Sudarshan Pant. 2023. A Review of Machine Learning and Deep Learning Approaches on Mental Health Diagnosis. *Healthcare* 11, 3 (2023), 285. https://doi.org/10.3390/healthcare11030285

[15] Kevin W. Jin, Qiwei Li, Yang Xie, and Guanghua Xiao. 2023. Artificial intelligence in mental healthcare: an overview and future perspectives. *British Journal of Radiology* 96, 1150 (2023), 20230213. https://doi.org/10.1259/bjr.20230213

[16] Dan W. Joyce, Andrey Kormilitzin, Katharine A. Smith, and Andrea Cipriani. 2023. Explainable artificial intelligence for mental health through transparency and interpretability for understandability. *npj Digital Medicine* 6 (2023), 6. https://doi.org/10.1038/s41746-023-00751-9

[17] Aziliz Le Glaz, Yannis Haralambous, Deok-Hee Kim-Dufor, Philippe Lenca, Romain Billot, Taylor C. Ryan, Jonathan Marsh, Jordan DeVylder, Michel Walter, Sofian Berrouiguet, and Christophe Lemey. 2021. Machine Learning and Natural Language Processing in Mental Health: Systematic Review. *Journal of Medical Internet Research* 23, 5 (2021), e15708. https://doi.org/10.2196/15708

[18] Han Li, Renwen Zhang, Yi-Chieh Lee, Robert E. Kraut, and David C. Mohr. 2023. Systematic review and meta-analysis of AI-based conversational agents for promoting mental health and well-being. *npj Digital Medicine* 6 (2023), 236. https://doi.org/10.1038/s41746-023-00979-5

[19] David D. Luxton. 2013. Artificial Intelligence in Psychological Practice: Current and Future Applications and Implications. *Professional Psychology: Research and Practice* 44, 6 (2013), 332–339. https://doi.org/10.1037/a0034559

[20] A. Luke MacNeill, Lillian MacNeill, Alison Luke, and Shelley Doucet. 2024. Health Professionals' Views on the Use of Conversational Agents for Health Care: Qualitative Descriptive Study. *Journal of Medical Internet Research* 26 (2024), e49387. https://doi.org/10.2196/49387

[21] Jared Moore, Declan Grabb, William Agnew, Kevin Klyman, Stevie Chancellor, Desmond C. Ong, and Nick Haber. 2025. Expressing stigma and inappropriate responses prevents LLMs from safely replacing mental health providers. *arXiv* (2025). https://doi.org/10.48550/arXiv.2504.18412 arXiv:2504.18412

[22] Aurelia Sauerbrei, Angeliki Kerasidou, Federica Lucivero, and Nina Hallowell. 2023. The impact of artificial intelligence on the person-centred, doctor-patient relationship: some problems and solutions. *BMC Medical Informatics and Decision Making* 23, 1 (2023), 73. https://doi.org/10.1186/s12911-023-02162-y

[23] Catherine E. A. Scipion, Margaret A. Manchester, Alex Federman, Yufei Wang, and Jalayne J. Arias. 2025. Barriers to and facilitators of clinician acceptance and use of artificial intelligence in healthcare settings: a scoping review. *BMJ Open* 15, 4 (2025), e092624. https://doi.org/10.1136/bmjopen-2024-092624

[24] Loujain Sharif, Reem Almabadi, Alhanouf Alahmari, Fai Alqurashi, Fidaa Alsahafi, Shahad Qusti, Walaa Akash, Alaa Mahsoon, Dev Bandhu Poudel, Khalid Sharif, and Rebecca Wright. 2025. Perceptions of mental health professionals towards artificial intelligence in mental healthcare: a cross-sectional study. *Frontiers in Psychiatry* 16 (2025), 1601456. https://doi.org/10.3389/fpsyt.2025.1601456

[25] Anna van Oosterzee. 2025. AI and mental health: evaluating supervised machine learning models trained on diagnostic classifications. *AI & Society* (2025). https://doi.org/10.1007/s00146-024-02012-z

[26] Mahima Anna Varghese, Poonam Sharma, and Maitreyee Patwardhan. 2024. Public Perception on Artificial Intelligence-Driven Mental Health Interventions: Survey Research. *JMIR Formative Research* 8 (2024), e64380. https://doi.org/10.2196/64380

[27] Liying Wang, Tanmay Bhanushali, Zhuoran Huang, Jingyi Yang, Sukriti Badami, and Lisa Hightow-Weidman. 2025. Evaluating Generative AI in Mental Health: Systematic Review of Capabilities and Limitations. *JMIR Mental Health* 12 (2025), e70014. https://doi.org/10.2196/70014

[28] Isabel B. Weir, Austin M. Stroud, Jeremiah J. Stout, Barbara A. Barry, Arjun P. Athreya, William V. Bobo, and Richard R. Sharp. 2025. Physician Perspectives on the Impact of Artificial Intelligence on the Therapeutic Relationship in Mental Health Care: Qualitative Study. *JMIR Mental Health* 12 (2025), e81970. https://doi.org/10.2196/81970

[29] Dong Whi Yoo, Sindhu Kiranmai Ernala, Bahador Saket, Domino Weir, Elizabeth Arenare, Asra F. Ali, Anna R. Van Meter, Michael L. Birnbaum, Gregory D. Abowd, and Munmun De Choudhury. 2021. Clinician Perspectives on Using Computational Mental Health Insights From Patients' Social Media Activities: Design and Qualitative Evaluation of a Prototype. *JMIR Mental Health* 8, 11 (2021), e25455. https://doi.org/10.2196/25455

[30] Melody Zhang, Jillian Scandiffio, Sarah Younus, Tharshini Jeyakumar, Inaara Karsan, Rebecca Charow, Mohammad Salhia, and David Wiljer. 2023. The Adoption of AI in Mental Health Care-Perspectives From Mental Health Professionals: Qualitative Descriptive Study. *JMIR Formative Research* 7 (2023), e47847. https://doi.org/10.2196/47847

[31] Tianlin Zhang, Ann-Marie Schoene, Shaoxiong Ji, and Sophia Ananiadou. 2022. Natural language processing applied to mental illness detection: a narrative review. *npj Digital Medicine* 5 (2022), 46. https://doi.org/10.1038/s41746-022-00589-7

[32] Andrea Zucchetti, Gabriele Nibbio, Luca Altieri, Lorenzo Bertorni, Irene Calzavara-Pinton, Elena Invernizzi, Nicola Necchini, Caterina Cerati, Laura Poddighe, Viola Bulgari, Jacopo Lisoni, Giacomo Deste, Stefano Barlati, and Antonio Vita. 2024. Artificial intelligence applications in mental health: the state of the art. *Italian Journal of Psychiatry* 10 (2024), 17–24. https://doi.org/10.36180/2421-4469-2024-5

A INTERVIEW GUIDE

Demographics

Age:
Gender:
Country of practice:
Primary professional role:
Highest completed education:
Years of experience:
Primary therapeutic orientation:
Type of organization (private practice, outpatient clinic, inpatient care, etc.):
Client population mostly treated (adults, adolescents, children, mixed):

Questions

Familiarity with AI

How familiar are you with AI technologies in general? SCORE

Diagnostic and assessment use

What are your thoughts on AI being used in diagnostic or assessment processes?

Monitoring and progress tracking

How do you view the role of AI in tracking patient progress over time?

Administrative and practical support

How do you feel about AI generating drafts of reports or assisting with scheduling and reminders?

Psychoeducation and between-session tools

What are your thoughts on using AI for psychoeducation or therapeutic exercises between sessions?

Risks and limitations

What risks or concerns do you have about using AI in clinical work?

Client perspective and therapeutic alliance

How do you think clients respond to AI being used in their treatment? SCORE

Professional identity, autonomy, and future conditions

Which aspects of your work should remain strictly human, and what would you need to feel confident about using AI safely in practice?