

**I:** The first question is about familiarity with AI. How familiar are you with AI technologies in general?

**P:** The classic ChatGPT and Microsoft Copilot I do use occasionally, but I try not to use them that much. There's not really any getting away from it nowadays though. Just the standard kind of use, I suppose.

**I:** Has your organization also introduced AI at all?

**P:** Not specifically. They did say they would prefer us to use Microsoft Copilot instead of ChatGPT because of the privacy thing, but they are still working on it. I saw a recent post on our intranet saying they want to start looking more into what it could offer, but they haven't done that yet.

**I:** On a scale of one to ten, how familiar do you think you are with AI in general?

**P:** I'd say a five. I know broadly what it does and how something like ChatGPT works, but the depths of it I don't really know.

**I:** Then for the next question, this is about diagnostics and assessment use. What are your thoughts on AI being used in diagnostic or assessment processes?

**P:** I would be skeptical about it. It depends on what kind of AI you use. From what I've been told, LLMs are more or less about estimating what the next word in a sentence could be, and I don't think that's a fit for diagnostics, because then it would just be based on what's most logical or most common online. So I wouldn't prefer using it for the actual research on the client, but I could imagine it being used for the administrative side, like reports afterwards, because that's a big hassle and takes a lot of time.

**I:** Do you also do diagnostics with your clients?

**P:** Yes.

**I:** Is there also a screening before the diagnostics start?

**P:** We usually do that. If I think there might be some trauma or ADHD, I start by letting them fill out a screener. It's in person or online, but I already know them usually by that point.

**I:** So that's based on information they already gave you. And how about summarizing patient information, do you think AI would be useful for that?

**P:** Yes and no. It could definitely be useful for making treatment plans or summarizing, but there's a big privacy issue that would need to be addressed. I would need to be sure it's safe to use, because I wouldn't want things to leak or be used for other AI purposes. I have a colleague who has a client that is really into AI and suggested they use it to summarize their session. They recorded it and ran it through AI, and it summarized really well, in much less time. If I have a

session of an hour I'm usually still typing for about 15 minutes trying to remember what we talked about, what we did, and what the homework was. I think it could be really convenient, but I would need to be sure it doesn't leak or is stored somewhere safe.

**I:** Then for the next question, this is about monitoring and progress tracking. How do you feel about the role of AI in tracking patient progress over time?

**P:** We use routine outcome management here, so we use questionnaires to measure complaints at the start, every few months, and at the end, to see whether they have gone down. I think AI could help with analysis of that, but on the other hand it's quite straightforward. We already have graphs showing whether things have gone up or down, so I don't think it's necessary.

**I:** How do you currently do the monitoring? At the beginning, the end, or also in the middle?

**P:** All of them. We do it at the beginning to see the baseline, then every three months, and depending on how long they are in treatment you can see different measurement points over time.

**I:** And are those measurements self-reported by the client or also in talks with you?

**P:** Both. Some are self-report, like the Forensic Complaints List, which covers things like anger issues and how clients experience things themselves. Then there are a few I fill in alone, like how the client's network or work is going. And there is one I usually fill out with the client, at least the first time, which is the risk evaluation to assess the risk of reoffending. We do that at the beginning and then every three to four months, though I don't always do it again with the client because by then I usually know how things are going and can fill it in myself.

**I:** So you generally don't think it would be that beneficial to use AI in this field?

**P:** No, I don't think so.

**I:** Then for the next question, this is about administrative and practical support. You briefly mentioned this already. How do you feel about AI generating drafts or reports or assisting with scheduling and reminders?

**P:** I think that could be really useful. A big part of our job is administration, even though that's not why I chose this profession. With diagnostics for example, it's usually a big research process to get to the final conclusions, which can take months. Then afterwards you have to look at all the tests, reports, everything, and still write it all down. That takes a lot of time, and you're also expected to see many clients. So if AI could support that it would be very useful, again only if privacy is accounted for. Writing reports of sessions could be helpful, although writing it myself also helps me re-evaluate the conversation and decide where to go next. But time-wise, it would save a lot.

**I:** How much time do you think you're actually spending on administrative tasks? Because I think the load is heavier at the end of treatment than during, right?

**P:** Yeah, definitely. At the end it's way more dense with administration. Just for an intake, you first call the client to plan it, then you read all the available information, you do the intake, then you write the report. You discuss it with your team whether or not to start treatment, and if you don't, you have to do the routine outcome management, make an end letter to the person or their referral, and make a closure form in the system. So just for one intake of an hour and a half, it's maybe four to six hours of work. And at the end of treatment as well, making sure everything is uploaded, that the end measurements are done, that the letters are there. It's more at the beginning and the end than in between.

**I:** Would you trust AI-generated reports or would you need to review everything carefully?

**P:** I would want to review it, especially with diagnostics. It's not just hard measurements like questionnaires, it's also your interpretation and your observations. It's not just what they filled in, it's also whether they understood the question, how they were in contact, what they did. So yes, I would want to review it.

**I:** What concerns would you have about accuracy or liability?

**P:** I don't really know whether AI is good enough to actually understand what it's doing. That may sound arrogant, but it can be complicated to integrate all of the information. I have tried to use it sometimes, anonymized of course, asking things like, I have these results, what does it mean? And it can help. But I don't know if there are other AI systems that work differently and are better focused on these kinds of situations. I would be skeptical. I think I could use it to get me going, but I would still use my own critical thinking to see if it's actually true and if it agrees with what I've seen.

**I:** You mentioned you tried it before, just to feel it out and test how it would perform?

**P:** Yeah. I used it for some diagnostic reports, asking what it means if I get certain results. I remember using it with an IQ test. I had the results and scored them myself, that wasn't the problem. But I loaded them into Microsoft Copilot to see what they could mean for someone in their day-to-day life, just to see if there was anything relevant for the client. It gave a nice overview, something like someone with this profile might lean towards this or show this behavior, and I thought that could be true for this person. I didn't use any names though.

**I:** Did it also give any wrong information or things that were clearly not relevant?

**P:** No, I think it was pretty spot on.

**I:** Then the next question is about psychoeducation and between-session tools. What are your thoughts on using AI for psychoeducation or therapeutic exercises between sessions?

**P:** On one hand, I don't really know. I've seen multiple articles about people using ChatGPT or other AI systems to make up for a therapist while they were on a waiting list, and I think that's really problematic and should not be done. There's no way for those people to check whether they are getting the right information. There are already great tools for psychoeducation, a lot online, and as a therapist you can usually help your client find them. At your workplace there are often already tools for that. So I would be careful with AI because you don't want clients to get wrong information. Even if it might not be wrong, it's not checked. You don't really see what they are getting.

**I:** What if the AI provides information that contradicts your treatment approach?

**P:** I think that could be problematic. People are really fond of AI and seem to take everything it says as the truth. I think it's dangerous if things are said that are not true. There are examples where people even got tips for suicide, asking how they should kill themselves and the AI just went along with it. That's very problematic, especially for vulnerable people or people prone to psychosis or schizophrenia. That's why I'm very hesitant to use it for that kind of thing.

**I:** Would you also worry about clients relying too much on AI between sessions?

**P:** For my clients specifically, I haven't heard any of them talk about using AI or ChatGPT, so I'm not really worried about that for them right now. It could be something I start asking about more. But in general, I do see an alarming use of AI for things that could actually be done with a certified psychologist. There are long waiting lists in mental health care and people resort to other options like AI, even though that's not great, but it is cheap and fast. I think that could be really harmful.

**I:** How would you monitor what the AI is telling your clients?

**P:** If I knew they were using it, I would be curious in the session and ask them if I could see the chat, what they learned from it, what they took away. Not in a degrading sense, just out of curiosity. And then I would try to warn them about the risks, especially if they are sharing very private things. I wouldn't blame them, but I would try to educate them.

**I:** Then the next question is about risks and limitations. What risks or concerns do you have about using AI in clinical work?

**P:** Privacy. Is it guaranteed, and can it actually be guaranteed? Also misinformation, especially towards clients, saying things that aren't true or aren't tailored to them. Tailoring is what we do in our work. And also what I mentioned about clients using it: the way AI tends to work around what someone says, it's always on the side of the user, which can feel very validating and empathetic. But my job as a therapist is also to hold up a mirror, to be empathetic but also to ask how things come across to other people. Especially if someone says dangerous things like planning to hurt themselves or others, that's a sign for me to take action, not to just validate it. So mainly the privacy risk and the risk of misinformation.

**I:** Have you thought about issues like bias?

**P:** That's partly what I mean with AI leaning toward what someone is already saying. It can be biased to support the user. And from my understanding it's based on what's already available on the internet, so if people with wrong intentions spread misinformation about depression or ADHD, that could also end up being used by my clients. I think it's biased in the sense that it's not based on actual facts but on what is most spread.

**I:** What concerns do you have about transparency or explainability of AI decisions?

**P:** I think that is a problem as well. Where does the information come from and how did it reach that conclusion? I noticed this when writing my thesis. I did use AI for support, like structuring thoughts or getting started, not to write it for me. But sometimes I asked for articles on certain topics and it would give full quotes and references to articles that didn't exist. It sounded very sure of itself, saying things like research shows this and that, but it turned out not to exist. So it's not really transparent. It presents itself as being based on research, but it's not.

**I:** Are there any other limitations you want to discuss?

**P:** There are probably a lot, but none that I can think of right now.

**I:** Then the next question is about client perspective and therapeutic alliance. How do you think clients respond to AI being used in their treatment?

**P:** I think that really differs between clients. From my experience, a lot of my clients are quite hesitant with technology and I don't think they would really like using AI. For example, when I want to record sessions for training or education, it's often difficult to get consent. So I think they wouldn't always want AI to be used either. But for more tech-savvy or younger clients who already use AI, it might help them get more engaged. So I think it really depends.

**I:** For education and training you sometimes want to record sessions and clients don't always agree. How often do you think that happens, percentage-wise?

**P:** That's difficult to say. I haven't asked that often, I really don't know.

**I:** On a scale from one to ten, how willing do you think clients are to use AI in their treatment?

**P:** I would say around a three or four. Some clients would be fine with it or interested, but most would be hesitant.

**I:** How might using AI affect the therapeutic alliance?

**P:** I don't really have experience with that, but I think it could be make or break. It could help with homework or staying engaged between sessions, since I see clients one hour a week and the rest happens at home. But it could also go wrong. If clients get different information from AI than from me, that could cause tension and make them more skeptical of my knowledge. That would be a difficult conversation.

**I:** Which aspects of your work should remain strictly human, and what would you need to feel confident about using AI safely in practice?

**P:** The sessions themselves should remain human. There is the option to do it online, but AI can't replace the human connection you feel when you're in the room together, especially with the risks I mentioned earlier. As for feeling confident, I would need some kind of guarantee that the data is handled safely and with care. I don't know how they would do that, but that's what I would need.

**I:** Why do you think the actual sessions need to stay human?

**P:** Mostly what I said earlier about the human critical ability. I doubt AI would be able to challenge clients or tell them to seek help if they are suicidal, to say something like I'm really worried, you need to see a doctor now. There needs to be some kind of crisis response involved. So that's my reason to keep the sessions human.

**I:** Do you have anything else to add?

**P:** I don't think so.