

I: Um, so the first question is an open question.

I: It's regarding familiarity with AI.

I: So how familiar are you with AI technologies in general?

P: Well, I mean the classic chatgpt and Microsoft Copilot and stuff I do use occasionally, but I try not to use it that much, but of course there's not really any getting away from it nowadays.

P: I think like, yeah.

P: So just, I guess, the standard kind of use of it, I know.

I: And has your organization also introduced AI at all?

P: Well, not specifically.

P: They did say that they would prefer us to use Microsoft Copilot instead of ChatGPT because of the privacy thing, I suppose.

P: But they haven't really, they are working on it.

P: I saw a recent post on our intranet that said that they want to start looking more into it, what it could offer us, but they haven't done that yet.

I: Okay.

I: Um, and for this question, we also want to know, on a scale of one to 10, how familiar you think you are with AI in general?

P: I'd say like a five.

P: Like I know kind of what it does and how it works, or how most AI like ChatGPT work, I think, broadly.

P: But the depths of it, I don't really know.

I: Okay.

I: Yeah.

I: Um, then for the next question, it's about diagnostics and assessment use. So what are your thoughts on AI being used in diagnostic or assessment processes?

P: Um, I would be kind of skeptical about it. It depends, I guess, what kind of AI you use. Like I know that there's the LLM, which kind of, like in my, what I've been told, is more or less about kind of estimating what the next word in the sentence could be.

P: And I don't think that will be a fit for diagnostics, because then it would just be based on what's the most logical or the most used online.

P: So I wouldn't prefer using it for the actual research on the client, but I could imagine it being used for the administrative side of things.

P: So like the reports afterwards, because that's a big hassle and it takes up a lot of time, which I think could be done quicker.

I: Okay, because you also do diagnostics with your clients or not?

P: Yes i do

I: And is there also a screening before the diagnostics start?

P: We usually do that, yeah. Like, you mean like regular screening for diagnostics, not for AI or something?

I: So before you have diagnostics, that's like an in-person conversation, right? Is there any screening, for example a questionnaire or some pre-research you do on the client?

P: Yes, but I do it like I have my own caseload.

P: And if I sort of think, well, maybe there might be some trauma here or ADHD, I start by letting them fill out a screener.

P: Yes, but it is in person or online, but still, like I already know them usually.

I: Yeah, so that's based on some information they already gave you.

I: And how about summarizing patient information? Do you think AI would be useful for that?

P: Yes and no. I think it could definitely be useful for making treatment plans or kind of summarizing it all, but there's a big privacy thing that would need to be provided for.

P: I would need to be sure that it's safe to use it, because I wouldn't want things to leak or be used for other AI stuff.

P: I think those are the kinds of things that AI could help with.

P: For example, I have a colleague who has a client who is really into AI.

P: The client suggested maybe we could use AI to summarize the session they had that day.

P: It was based on the client's suggestion.

P: They did record it all and let it run through AI.

P: They summarized it really well.

P: It was a very good summary of the session and it took way less time.

P: Like, for example, if I had a session of an hour, I'm usually still typing for like 15 minutes trying to remember what we talked about, what we did, and what the homework was.

P: I think it could be really convenient, but I would need to be sure that it doesn't leak in any way or is stored somewhere safe.

I: Yeah.

I: Makes sense.

I: Okay, then for the next question, it's about monitoring and progress tracking.

I: How do you feel about the role of AI in tracking patient progress over time?

P: Well, we use routine outcome management here.

P: So we use questionnaires to measure the “klachten” —*complaints*— at the start and every few months and at the end, to see whether or not it has gone down.

P: I think AI could help with analysis of that, but on the other hand it's quite straightforward.

P: We already have graphs showing whether it's gone up or down, so I don't think it's necessary.

I: So how do you currently do the monitoring?

I: Do you do it at the beginning, at the end, or also in the middle?

P: Yeah, all of them actually.

P: We do it at the beginning to see what's kind of the baseline.

P: And then we do it every three months, I would say.

P: And then depending on how long they are in treatment, like say it's a year, you can see different measurement points.

P: And then you can see whether or not it's gone down.

I: And those measurements, are those self-reported by the client or also in talks with you?

P: Both actually.

P: Some are self-report, like for example we use the Forensische Klachtenlijst —*Forensic complaints list*—.

P: That's about things like anger issues or how they experience it themselves.

P: Then we have a few that I alone fill in, like how the client's network is going or how work is going.

P: And then we have one that I usually fill out with the client, at least the first one, because we need information I don't have if I don't know them yet.

P: It's the risk evaluation, to figure out the risk of reoffending.

P: We do that at the beginning and then every three to four months as well.

P: But I don't always do it again with the client, because I usually know how things are going, so then I fill it in myself.

I: Yeah.

I: So you generally don't think it would be that beneficial to use AI in this field?

P: No, I don't think so.

I: Okay.

I: Then for the next question, this is about administrative and practical support.

I: You briefly mentioned this already.

I: So how do you feel about AI generating drafts or reports or assisting with scheduling and reminders?

P: Yeah, I think that could be really useful.

P: Do you mean reminders for the clients?

I: Both you and the client, these questions are open to explore where things go, so whatever you think we should talk about is valid here.

P: Well, in general, like I said, a big part of our job is administration, even though that's not why I chose this profession.

P: It's not the most fun part.

P: I don't really mind it, I like writing and reporting.

P: But for example with diagnostics, it's usually a big research process to get to the final conclusions, which can take months.

P: Then afterwards you have to look at all the tests, reports, everything, and then still write it all down.

P: That takes a lot of time, and you're also expected to see many clients.

P: You don't always have the time for that.

P: So if AI could support that, it would be very useful.

P: Again, only if privacy is accounted for.

P: Writing reports of sessions could be helpful.

P: Although writing it myself also helps me re-evaluate the conversation and decide where to go next.

P: But time-wise, it would save a lot of time if that could be done with AI.

I: Yeah.

I: So how much time do you think you're actually spending on administrative tasks?

P: On a day-to-day basis?

I: I mean, in general. Because I don't know if the load is way higher, like as you mentioned, I think in the end it's more administrative tasks than during the whole treatment, right?

P: Oh yeah, you mean at the end of a treatment?

P: Yeah, definitely.

P: At the end it's way more dense with the administrative side of things.

P: Just for what is called an intake, you first call the client to plan the intake, then you have to read all the stuff that is available, you do the intake, then you write the report of the intake.

P: You discuss it in your team whether or not you want to start this treatment.

P: And if we don't, then we have to do the routine outcome management thing.

P: And then we have to make an end letter to the person or their referral.

P: And then we have to make a closure form in our system.

P: So just for one intake of one and a half hours, it's like maybe four to six hours of work.

P: And then at the end as well, making sure everything is uploaded into our system, that the end measurements are done, that the letters are there.

P: So yeah, you're right about it.

P: It's more at the beginning and the end than in between.

I: Yeah.

I: Makes sense.

I: Um, so would you trust AI-generated reports or would you need to review everything carefully?

P: I would want to review it.

P: Yeah.

I: Yeah.

P: Especially with diagnostics and stuff.

P: It's not just hard measurements like questionnaires.

P: It's also your interpretation and your observations.

P: So it's not just they filled in something.

P: It's also did they understand the question, how are they in contact, what did they do.

P: So yeah, I would want to review it, definitely.

I: So what concerns would you have about accuracy or liability?

P: Well, like I mentioned earlier, I don't really know whether or not AI is good enough to actually understand what it's doing.

P: That may sound arrogant, but it can be complicated to integrate all of the information.

P: I have tried to use it sometimes, anonymized of course.

P: Like, I have these results, what does it mean?

P: And it can help.

P: But then again, I don't know if there are other AI systems that work in a different way that are better focused on these kinds of situations.

P: I would be skeptical.

P: I think I could use it to get me going, but I would still use my own critical thinking to see if it's actually true and if it agrees with what I've seen.

I: So you mentioned you tried it before.

I: That was just to feel it out and test how it would perform?

P: Yeah, kind of.

P: I used it for some diagnostic reports, like what does it mean if I get this and this.

P: I remember I used it with an IQ test.

P: I had the results and scored them myself.

P: That wasn't the problem.

P: But I loaded them into Microsoft Copilot to see what this could mean for someone in their day-to-day life.

P: Just to see if there was anything suitable for the client.

P: It gave a nice overview.

P: Like someone with this profile might lean towards this or show this behavior.

P: And I thought that could be true for this person.

P: So I did use it.

P: But I didn't use any names or stuff.

I: Yeah.

I: And did it also give any wrong information or things that were clearly not relevant?

P: No, I think it was pretty spot on.

I: Okay.

I: Then I think we can move on to the next question.

I: This is about psychoeducation and between-session tools.

I: What are your thoughts on using AI for psychoeducation or therapeutic exercises between sessions?

P: On one hand, I don't really know.

P: I've seen multiple articles about people using ChatGPT or other AI systems to kind of make up for a therapist while they were on a waiting list.

P: And I think that's really problematic and should not be done.

P: There's no way for those people to check whether they are getting the right information.

P: I think that goes for this as well.

P: There are already great tools for psychoeducation.

P: There's a lot online.

P: And as a therapist, you can usually help your client find them.

P: At your workplace there are often already tools for that.

P: So I would be careful with AI because you don't want clients to get wrong information.

P: Even if it might not be wrong, it's not checked.

P: You don't really see what they are getting.

I: So what if the AI provides information that contradicts your treatment approach?

P: I think that could be problematic.

P: People are really fond of AI and seem to take everything it says as the truth.

P: Not that I want to say that I know all the truth.

P: But I think it's dangerous if things are said that are not true.

P: You've probably seen examples where people even got tips for suicide.

P: Like asking how should I kill myself, and it just went along with it.

P: I think that's very problematic.

P: Especially for vulnerable people or people prone to psychosis or schizophrenia.

P: It's really dangerous.

P: That's why I'm very hesitant to use it for that kind of thing.

I: Yeah.

I: So would you also worry about clients relying too much on AI between sessions?

P: For my clients specifically, or in general?

I: Your interpretation

P: Well, I haven't heard any of my clients talk about AI or using ChatGPT.

P: So I'm not really worried about that for them right now.

P: It could be something that I could start asking more, actually.

P: But in general, yeah, I do see a “zorgwekkende” —*alarming*— use of AI for things that could actually be done with a certified psychologist.

P: I think that deals with a big problem in mental health care in general, which is that there are long waiting lists.

P: People resort to other options like AI, even though that's not great, but it is cheap and fast.

P: I think that could be really harmful.

I: Because how would you monitor what the AI is telling your clients?

P: If I knew that they were using it, I think I would be curious in the session and ask them, like, can I see the chat?

P: What did you learn from it?

P: What did you take away?

P: Not in a degrading sense.

P: I would just try to be curious and see what it said.

P: And then I would try to warn them about the risks of using AI for that.

I: Yeah.

P: Especially if they are telling very private stuff as well.

I: Yeah.

P: I wouldn't blame them, but I would try to educate them.

I: Yes.

I: Yeah.

I: Um, okay.

I: So then we can go to the next question.

I: This is about risks and limitations.

I: What risks or concerns do you have about using AI in clinical work?

P: Like I said, the privacy thing.

P: Is it guaranteed or not?

P: And can they actually guarantee it?

P: Is that even possible?

P: Also misinformation, especially towards clients, saying things that just aren't true or aren't tailored to them.

P: That's what we do in our work.

P: We tailor it to our clients to see what works for them.

P: And also what you said about clients using it, the way AI tends to work around what someone says.

P: It's always kind of on the side of the user.

I: Yeah.

P: Which can feel very validating and empathetic.

P: But my job as a therapist is also to hold up a mirror.

P: To be empathetic, but also to ask how things come across to other people.

P: Especially if someone says dangerous things like planning to hurt themselves or others.

P: That's a sign for me to take action, not to just validate it.

P: So mainly the privacy risk and the risk of misinformation.

I: Yeah.

I: And have you thought about issues like bias?

P: You mean that AI is biased toward some -?

I: Yeah.

P: Yeah.

P: That's partly what I mean with leading toward what someone is already saying.

P: It can be biased to support the user.

P: And from my understanding, it's based on what's already available on the internet.

P: So if people with wrong intentions spread misinformation about depression or ADHD, for example, that could also end up being used by my clients.

P: So I think it's biased in the sense that it's not based on actual facts, but on what is most spread.

I: Yeah.

I: And what concerns do you have about transparency or explainability of AI decisions?

P: I think that is a problem as well.

P: Like, where does the information come from and how did it reach that conclusion?

P: I noticed this when writing my thesis.

P: I did use AI for support, like structuring thoughts or getting started.

P: Not to write it for me.

P: But sometimes I asked for articles on certain topics.

P: And it would give full quotes and references to articles that didn't exist.

P: It sounded very sure of itself, like research shows this and that.

P: But then it turned out not to exist.

P: So it's not really transparent.

P: It presents itself as being based on research, but it's not.

I: Yeah, makes sense.

I: Are there any other limitations you want to discuss?

P: There are probably a lot, but not really any that I can think of right now.

I: Okay.

I: Then the next question is about the client perspective and therapeutic alliance.

I: How do you think clients respond to AI being used in their treatment?

P: I think that really differs between clients.

P: From my experience, a lot of my clients are quite hesitant with technology.

P: I don't think they would really like using AI.

P: For example, when I want to record sessions for training or education, it's often difficult to get consent.

P: Not everyone wants that.

P: So I think they wouldn't always want AI to be used either.

P: But for more tech-savvy or younger clients who already use AI, it might help them get more engaged.

P: It could attract them.

P: So I think it really depends.

I: Yeah.

I: So for education and training, you sometimes want to record sessions, and clients don't always agree.

I: How often do you think that happens, percentage-wise?

P: That's difficult to say.

P: I haven't asked that often.

P: I really don't know.

I: That's okay.

I: For this question, we also wanted a score.

I: On a scale from one to ten, how willing do you think clients are to use AI in their treatment?

P: I would say around a three or four.

P: Some clients would be fine with it or interested, but most would be hesitant.

I: Yeah.

I: How might using AI affect the therapeutic alliance?

P: That's a good question.

P: I don't really have experience with that.

P: But I think it could be make or break.

P: It could help with homework or staying engaged between sessions.

P: I see clients one hour a week, and the rest happens at home.

P: But it could also go wrong.

P: If clients get different information from AI than from me, that could cause tension.

P: It could make them more skeptical of my knowledge.

P: That would be a difficult conversation, like how do we handle that, if they question my knowledge.

I: Okay, I think this is the last question.

I: Which aspects of your work should remain strictly human, and what would you need to feel confident about using AI safely in practice?

P: The sessions themselves should remain human, ofcourse there is option to do it online, but

P: AI can't replace the human connection you feel when you're in the room together.

P: Especially with the risks I mentioned earlier.

P: And what was the second part of the question?

I: What would you need to feel confident about using AI safely in practice?

P: I would need some kind of guarantee that the data is handled safely and with care.

P: I don't know how they would do that, but that's what I would need.

I: Mm-hmm.

I: Um, yeah.

I: So why do you think these aspects need to stay human?

I: Like the actual sessions?

P: Yeah.

P: I think mostly what I said earlier about the human critical ability.

P: I doubt if AI would be able to inhibit clients or tell them to seek actual help.

P: Or, well, actual help sounds degrading, but I mean like if they are suicidal.

P: Would an AI eventually say you need to go to a doctor now, I'm really worried?

P: There needs to be some kind of crisis thing involved.

P: I don't know.

P: So that's my reason to keep the sessions human.

I: Yeah.

I: Um, okay.

I: That's it.

I: So do you have anything else to add or to retract?

P: I don't think so.

I: Nope.

I: Okay.

I: Then I will stop the recording.